
Introduction to Algorithm Engineering

Spring 2022

Lecture 2

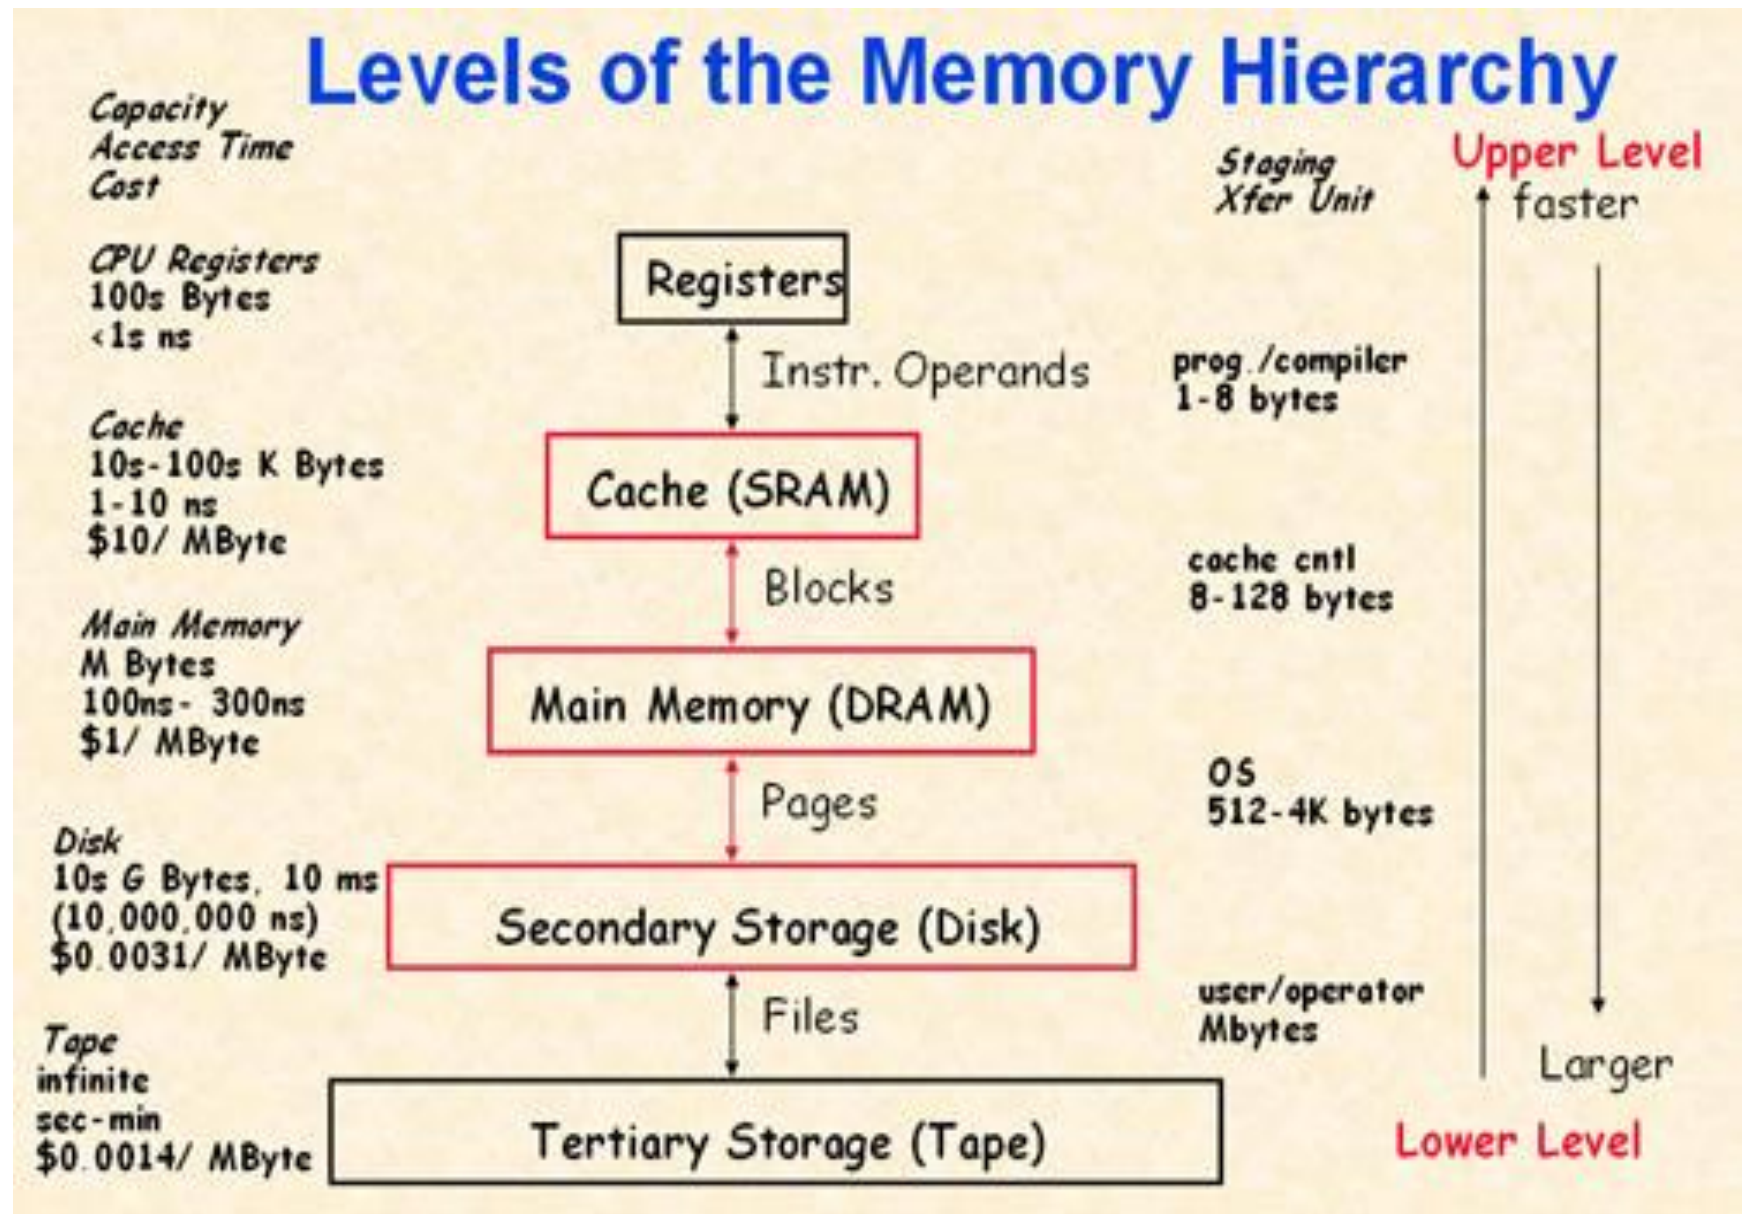
This Lecture

- In this lecture, we will study algorithm implementation from an architecture friendliness point of view.
- The central question to answer is: **How Far Should Algorithms Consider the Architecture?**
- There are predominantly two schools of thought
 - Very much
 - Should never
- Both sides have good arguments.

Cache Aware Algorithm Design

- A model that is relevant to efficient implementations of algorithms.
- Models the twin challenges of algorithm design for big data and a deep memory hierarchy.
- For instance, consider the case that the entire data does not fit in the memory of a computer.
- Then, the assumption that all memory references require the same access time is not truly valid.
- How to design and implement algorithms when the above assumption does not hold?

Example Hierarchy



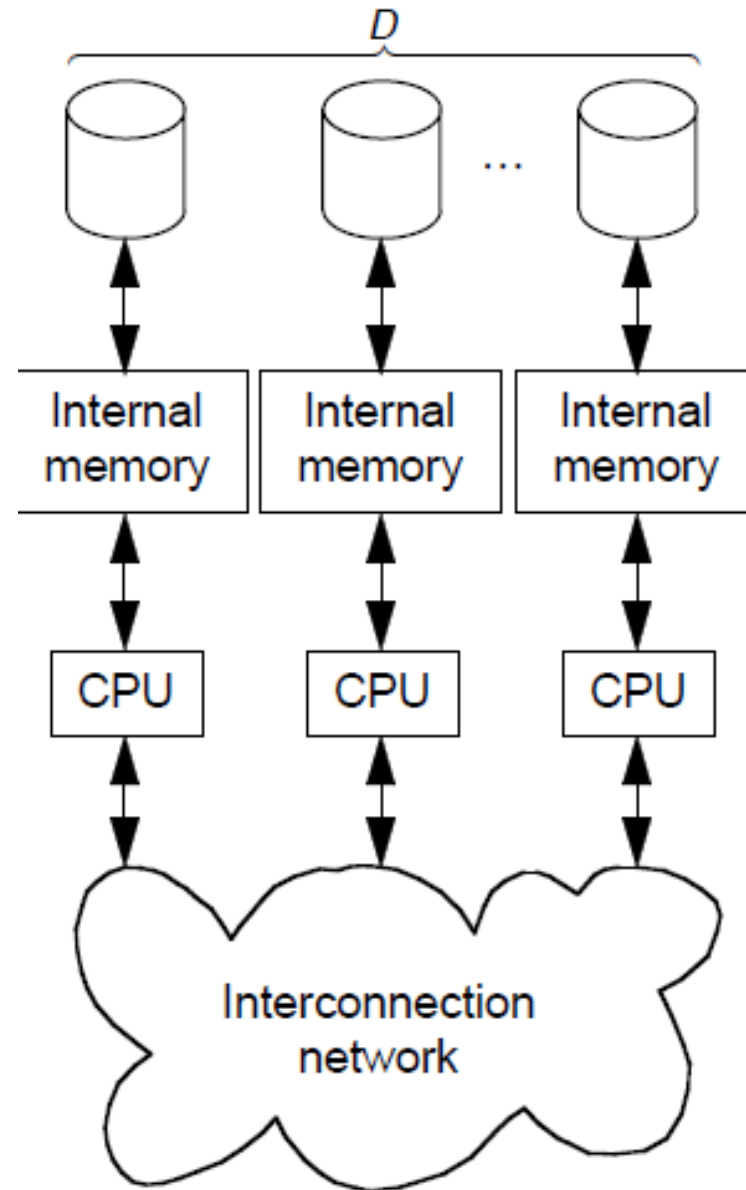

Some Definitions

- **Working Set:** A collection of data items recently accessed by a program.
- A small working set is good. Can fit in the cache, and a prefetcher can help in most cases minimize cache misses.
- However, prefetch works best for computations with very strong locality in the computation.
 - Example: Add elements of two arrays index-wise.
- But, there are other computations that have strong non-local access patterns.
- In such cases, explicit memory management helps.

Performance Parameters

- » The number of I/O operations,
 - » The disk space usage, and
 - » The CPU time.
-
- » N = problem size (in units of data items);
 - » M = internal memory size (in units of data items);
 - » B = block transfer size (in units of data items);
 - » D = number of independent disk drives;
 - » P = number of CPUs
-
- » Notice that $M < N$, and $1 \leq DB \leq M/2$

Model in Pictures



Four Fundamental Operations

- » **Scanning** (a.k.a. streaming or touching) a file of N data items, which involves reading or writing the items in the file sequentially.
- » **Sorting** a file of N data items, which puts the items into sorted order.
- » **Searching** online through N sorted data items.
- » The fourth one, is on **batched query** style operations. Read from Vitter's survey.

Scanning

- » The I/O bound $\text{Scan}(N) = O(N/B)$.
- » Relatively straight-forward to understand.
- » Sorting requires some additional tricks.

Sorting

- » We will show that sorting N items with a buffer ratio B requires $O(N/B \log_{(M/B)} N/B)$ disk I/O operations.
- » Two main techniques
 - ♦ Sample (Distribution) sort
 - ♦ Sorting by merging

Sample Sort

- » First the basic idea
- » Imagine an extension of quick sort.
- » Choose $s - 1$ pivots. (For quick sort, $s = 2$).
- » Use these pivots to divide the elements into s buckets.
- » At this point, elements in bucket i are all smaller than those in bucket $i+1$, for $i=1,2,\dots, s-1$.
- » Sort each bucket recursively.

Example Sample Sort

Sample Sort

- » Just as in quick sort, need good pivots so that each bucket has roughly the same number of elements.
- » Suppose we can find such pivots.
- » Let us analyze the number of I/O operations.

Sample Sort

- » For equal sized buckets, at every recursive level, the size of the bucket drops by $O(s)$.
- » So, there are $O(\log_s N/M)$ levels of recursion.
 - ♦ Why N/M ?
- » In each level of the recursion, we have to scan the entire data across all the subproblems.
- » At $s = O(M/B)$, we get the required bound on the number of I/Os.

How to Find Good Pivots Quickly?

- » Several techniques exist.
- » One is find the k th quantiles.
 - ♦ Guaranteed to be good
- » Randomized methods also are possible.
- » More techniques for the multiple disk model.
- » Read from the survey of Vitter.

Sample Sort – Multiple Disks

- With multiple disks, the algorithm gets more complicated.
- Need to ensure that buckets are written into all disks in near equal quantities.
 - Otherwise, cannot use multiple reads at the same time.