

Merge based Sorting

- » Very similar to the standard merge sort technique.
- » Phase I: Called the **Run-Formation** Phase
 - ♦ Scan one buffer of data of size B
 - ♦ Sort the buffer internally
 - ♦ Write back the sorted buffer
- » Phase II: **Merge** Phase
 - ♦ Scan two buffers of data
 - ♦ Merge them
- » In both phases, one can use double buffering to overlap computation with disk I/O.

Transfer $B^1 B^2 B^3$
compute $\frac{B^1 B^2}{B^1 B^2}$

Merge Based Sorting

- So, the number of passes is $O(\log_{M/B} N/B)$.
- Using software prefetching can help sorting algorithms.
- As an aside, forming a permutation of n numbers is a special case of sorting.
 - So, similar bounds apply to permuting.

Matrix Multiplication

- A good case study of how to tune algorithms to a cache size.

```
MatrixMultiply(A, B, C, n)
Begin
  for i = 1 to n do
    for j = 1 to n do
      for k = 1 to n do
        C[i,j] += A[i,k]*B[k,j]
      end-for
    end-for
  end-for
End.
```

Matrix Multiplication

- A good case study of how to tune algorithms to a cache size.

```
MatrixMultiply(A, B, C, n)
Begin
  for i = 1 to n do
    for j = 1 to n do
      for k = 1 to n do
        C[i,j] += A[i,k]*B[k,j]
      end-for
    end-for
  end-for
End.
```

- Suffers from lots of cache misses.
- Exercise: Estimate the number of cache misses.

Matrix Multiplication – Naive




- To compute each element of C:
 - Read/Scan the entire row of A – $O(N/B)$ cache misses
 - Read/Scan the column of B – $O(N)$ cache misses.
 - Why?
 - Cache misses per element of C = $O(N) + O(N/B) = O(N)$.
- Overall cache misses = $O(N^3)$

Matrix Multiplication

- It is however better to imagine that the computation be broken into blocks of $s \times s$ submatrices.

```
MatrixMultiply(A, B, C, n)
Begin
  for i = 1 to n/s do
    for j = 1 to n/s do
      for k = 1 to n/s do
        C = C + Mul(Aik, Bkj, s)
      end-for
    end-for
  end-for
End.
```

A




j $j+s$




i




j



B



Big




Cache miss \approx page fault



Small
= cache


Matrix Multiplication in Pictures



$$C_{ij} = A[i:i+s, 1:1+s] \times B[1:1+s, j:j+s] +$$

$$A[i+s+1:i+2s, 1:1+s] \times B[1:1+s, j+s+1:j+2s] +$$

Improved Version



- Estimate the number of cache misses again.

for $i = 1$ to $\frac{n}{s}$ do

for $j = 1$ to $\frac{n}{s}$ do

for $k = 1$ to $\frac{n}{s}$ do

$C_{ij} += \text{Mul}(A_{ik}, B_{kj}, s) // s^3 \text{ comp}$



for one C_{ij} #cache misses

$$\#C_{ij} = \left(\frac{n}{s}\right)^3$$

$$= O\left(\frac{s^2}{B}\right) \left\{ \begin{array}{l} \text{Total} = \left(\frac{n}{s}\right)^3 \times \frac{s^2}{B} \\ = O\left(\frac{n^3}{sB}\right) \end{array} \right.$$

Improved Version

- Estimate the number of cache misses again.
- To compute an $s \times s$ submatrix of C , we need to bring in $O(s^2)$ elements into the cache.
- Results in $O(s^2/B)$ cache misses.

Improved Version

- Estimate the number of cache misses again.
- To compute an $s \times s$ submatrix of C , we need to bring in $O(s^2)$ elements into the cache.
- Results in $O(s^2/B)$ cache misses.
- For the entire C matrix, the number of cache misses = $O(s^2/B) \times (n/s)^3 = O(n^3/sB)$.
 - Why?

iterations

Improved Version

- Estimate the number of cache misses again.
- To compute an $s \times s$ submatrix of C , we need to bring in $O(s^2)$ elements into the cache.
- Results in $O(s^2/B)$ cache misses.
- For the entire C matrix, the number of cache misses = $O(s^2/B) \times (n/s)^3 = O(n^3/sB)$.
 - Why?
 - The program has a three way nested loop of n/s iterations in each loop.



Improved Version

$$s \left[\begin{array}{c} s \\ B \end{array} \right] \rightarrow s/B$$

- Estimate the number of cache misses again.
- To compute an $s \times s$ submatrix of C , we need to bring in $O(s^2)$ elements into the cache.
- Results in $O(s^2/B)$ cache misses.



- For the entire C matrix, the number of cache misses = $O(s^2/B) \times (n/s)^3 = O(n^3/sB)$.



good choice of s

- Tall cache Assumption: $M = \Omega(B^2)$, i.e., many more cache lines fit in the cache.

- Number of cache misses = $O(n^3/BM^{1/2})$ at $s = B$.

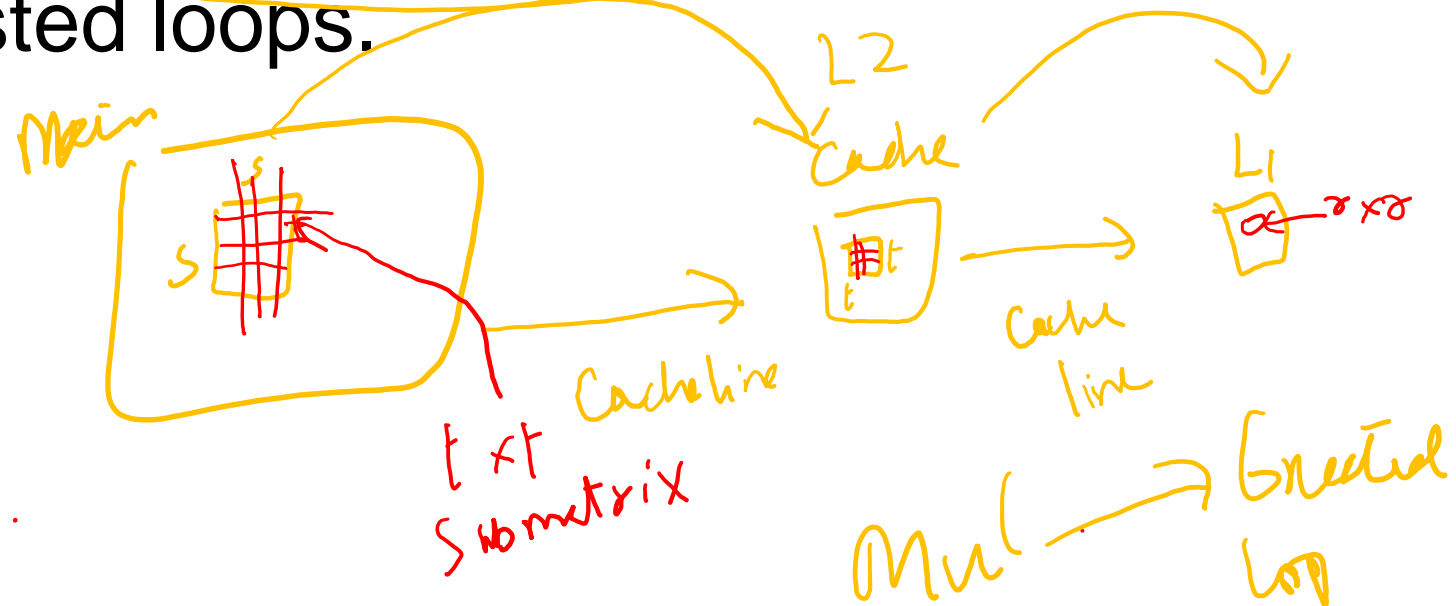
Improved Version

- Estimate the number of cache misses again.
- To compute an $s \times s$ submatrix of C , we need to bring in $O(s^2)$ elements into the cache.
- Results in $O(s^2/B)$ cache misses.
- For the entire C matrix, the number of cache misses = $O(s^2/B) \times (n/s)^3 = O(n^3/sB)$.
- Tall cache Assumption: $M = \Omega(B^2)$, i.e., many more cache lines fit in the cache.
- Number of cache misses = $O(n^3/BM^{1/2})$ at $s = B$.

Further Tuning

- May be required if there are more cache levels.
- Each $s \times s$ submatrix has to be further blocked for the next level of the memory.
 - Number of such blocking parameters increases with the number of levels in the memory hierarchy.
- Increases the program complexity because of multiple nested loops.

Intel MKL
- Subroutine for
matrix multiply



Summary

- Blocking is a good technique for making an implementation cache-aware.
- There are some lower bounds too.